University of California, Berkeley
College of Engineering
Computer Science Division – EECS

Spring 2018   Anthony D. Joseph and Jonathan Ragan-Kelley

Midterm Exam #3
April 25, 2018
CS162 Operating Systems

Your Name:	
SID AND 162 Login:	
TA Name:	
Discussion Section Time:	

General Information:
This is a **closed book and TWO 2-sided handwritten notes** examination. You have 110 minutes to answer as many questions as possible. The number in parentheses at the beginning of each question indicates the number of points for that question. You should read **all** of the questions before starting the exam, as some of the questions are substantially more time consuming.

Write all of your answers directly on this paper. *Make your answers as concise as possible.* If there is something in a question that you believe is open to interpretation, then please ask us about it!

Good Luck!!

QUESTION	POINTS ASSIGNED	POINTS OBTAINED
1	36	
2	22	
3	22	
4	20	
TOTAL	100	

NAME: _____

1. (30 points total) Short answer
   a. (12 points) True/False and Why? **CHECK THE BOX FOR YOUR ANSWER.**
      i) Pintos uses separate user-kernel mode transfer implementations for interrupts and syscalls.

❏ TRUE   ❏ FALSE

Why?

      ii) One only needs a source & destination IP address, source & destination port to uniquely identify a network connection.

❏ TRUE   ❏ FALSE

Why?

      iii) The Two Phase Commit protocol guarantees that all the participants commit or abort at the same time.

❏ TRUE   ❏ FALSE

Why?

      iv) Applying a conservative view of the End-to-End Principle to reliable file transfer means you should only implement reliability functionality at the endpoints.

❏ TRUE   ❏ FALSE

Why?

NAME: _____

b. (14 points) Consider the following 6 I/O operations and their respective cylinder locations on disk. Seek time is 0.1milliseconds *per cylinder traversed*.

Operation	A	B	C	D	E	F
Cylinder	4	10	35	62	69	95

i) Use the Shortest Seek Time First disk scheduling algorithm to schedule these operations. The arm begins at cylinder 33. List the order in which the 6 operations are scheduled and calculate the total seek time. *Ignore rotational and transfer delays.*

Letter:

_____, _____,_____,_____,_____,_____

Total seek time: _____

ii) Unhappy with this seek time, we decide to use SCAN as our disk scheduling algorithm instead. Assume SCAN begins by traversing descending cylinder numbers starting with the disk arm at cylinder 33. List the order in which our 6 operations are scheduled and calculate the total time we spend seeking. You may ignore rotational delay in your calculations.

Letter:

_____, _____,_____,_____,_____,_____

Total seek time: _____

c. (3 points) Journaling is a technique that can be used to ensure that a file system recovers its metadata to a consistent state after a crash. One type of journaling is called "**data journaling**," because all file system metadata and data is written to the log before being updated in place. Describe the performance of such a journaling file system, as compared to a file system that only logs metadata, for *very large* sequential write workloads. *Roughly* what throughput can you expect and why?

NAME: _____

d. (4 points) What do the four letters in ACID stand for?

A_____

C_____

I_____

D_____

e. (3 points) <u>Briefly</u> explain how from a programmer's prospective, using RPC differs from using procedure calls.

NAME: _____

2. (22 points total) Revisiting Homework #2.

Your new startup lets users watch streaming anime videos, so you reuse the code you wrote for homework #2 to create a simple video server that accepts client connections, reads in the video title string (e.g. "Kimi no Na wa."), and writes back the raw HD video file to the client.

- Ignore HTTP formatting and headers – assume we read & write raw strings.
- Make sure your code is multithreaded and allows handling of multiple simultaneous client connections.
- Assume very large video files (possibly on the order of gigabytes or more).
- Do not forget to clean up and not leak resources, including freeing memory and closing file descriptors when appropriate.

a) (16 points) Fill in the below code, <u>including all blank lines and freeform boxes</u>. Assume for this part that system calls will succeed when possible (i.e. you do not have to include, for example, checks for thread creation or malloc success). Also assume that the client connection does not drop or hang up prematurely.

```
#define MAX_LEN 20
#define BUFFER_SIZE 4096

void *serve_video(void *);

int main() {
  pthread_t thread; struct sockaddr serv_addr,
                              cli_addr; int cli_len;
  // sockaddr initialization code omitted
  int sockfd = _____(PF_INET, SOCK_STREAM, 0);
  _____(sockfd, (struct sockaddr *) &serv_addr,
                              sizeof(serv_addr));
  _____(sockfd, 5);
  while (1) {
      int newsockfd = _____(sockfd, (struct
                  sockaddr *) &cli_addr, &cli_len);
      _____
  }
  _____
  return 0;
}
```

NAME: _____

```
void *serve_video(void *arg) {
  int newsockfd = (int) arg;
  // expect both "request" and "path" < 20 chars
  char request[MAX_LEN];
  char path[MAX_LEN];
  int bytes_read = read(newsockfd, request, MAX_LEN);
  int success = request_to_video_path(request, path);
  // request_to_video_path(): takes "request" string,
  // populates "path" string with file path to video
  if (success == -1) { //request not found or invalid

    _____
    return NULL;
  }
  int fp = open(path, O_RDONLY);
  char* buffer = malloc(BUFFER_SIZE*sizeof(char));
  /* Fill in the following box */
```

┌──┐
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
│   │
└──┘

```
  return NULL;
}
```

NAME: _____

b) (3 points) Assume you coded part a) correctly. Now suppose that malloc() failed in one the threads running `serve_video()`. Briefly, in 1-2 sentences, explain what would happen and why.

c) (3 points) Suppose you forgot to close any of your file descriptors in the above code and you run it indefinitely. Briefly, in 1-2 sentences, explain what could happen to your video server and why.

NAME: _____

3. (22 points total) Filesystems.

The Berkeley Fast File System (FFS) in BSD 4.2 was one of the first file systems that treated the disk like a disk and thus improved performance through spatial locality

   a. (3 points) What on-disk structures does FFS use to track allocation of inodes and data blocks?

   b. (3 points) Why was the innovation in (3a) an improvement over the prior Unix file system in BSD 4.1?

   c. (4 points) Recall that FFS tries to spread large files across disk by splitting them into chunks and putting each chunk in a different part of a block group. Given a hard drive with a peak transfer rate of 200Megabytes/s and that seek and rotation combined take on average 16 milliseconds. How big should each chunk of the file be so that we can achieve 2/3 of the peak transfer rate for large files when they are accessed sequentially? *Show work outside of the box for partial credit.*

   d. (12 points) List the set of disk blocks that must be read into memory in order to read the file /cs162/bar/foo.txt in its entirely from a UNIX BSD 4.2 filesystem. Assume that the file is 12,300 Bytes long and that the disk blocks are 1 Kilobytes long. You can assume that each directory fits into a single disk block.

   Number of Blocks:

   Blocks:

NAME: _____

4. (20 points total) Security and Two-Phase Commit.

Suppose that No Such Agency headquarters (coordinator) in **Washington** (W) communicates with branch offices (workers) in **Moscow** (M) and **Beijing** (B) using symmetric-key-encrypted, reliable messages. Suppose furthermore that the encryption key used is changed each day; every evening the new encryption key that is to be used the next day is sent from Washington to the two branch offices. Two-Phase Commit is used in order to ensure that, despite crashes, either (a) everyone eventually switches to the new encryption key, NEWKEY, or (b) no one switches to the new encryption key and everyone continues using OLDKEY. Assume that a machine takes a *long* time to reboot and recover after a crash, and timeouts are 1 second.

The steps involved in implementing two-phased commit are listed below, in time order:
1. Washington: write "begin transaction" to its log
2. Washington → Moscow: "New key is NEWKEY." (→ means W sends msg to M)
3. M: write "New key is NEWKEY." to its log
4. M → W: "Prepared to commit"
5. W → B: "New key is NEWKEY."
6. B: write "New key is NEWKEY." to its log
7. B → W: "Prepared to commit"
8. W: write "New key is NEWKEY." to its log
9. W: write "commit" to its log
10. W → M: "commit"
11. M: write "got commit" to its log
12. M: Key = NEWKEY
13. M → W: "ok"
14. W → B: "commit"
15. B: write "got commit" to its log
16. B: Key = NEWKEY
17. B → W: "ok"
18. W: Key = NEWKEY
19. W: write "done" to its log

    a. (4 points) If W crashes after step 9 and no one else fails, what key will everyone end up using, once W reboots and recovers? ***Give the reason why.***

       Key: _____

       Why:

NAME: _____

    b. (4 points) If M crashes after step 4 and no one else fails, what key will everyone end up using, once M reboots and recovers? ***Give the reason why.***

> **Key:** _____

> **Why:**

    c. (4 points) If B crashes after step 4 and no one else fails, what key will everyone end up using, once B reboots and recovers? ***Give the reason why.***

> **Key:** _____

> **Why:**

NAME: _____

d. (4 points) If M crashes after step 10, what recovery steps must it take after it reboots in order to achieve the correct global state with respect to which encryption key to use?

```



```

e. (4 points) If W crashes after step 11, what recovery steps must it take after it reboots in order to achieve the correct global state with respect to which encryption key to use?

```



```

NAME: _____

5. (0 points total) That's All, Folks!

    a. What does the "r" in "jrk" stand for? (You should really know this one)

    b. Favorite part of CS162?

    c. Any other thoughts on how CS162 staff is doing or to improve the course?

Hope you enjoyed CS162 and operating systems,
have a fun summer!