University of California, Berkeley
College of Engineering
Computer Science Division — EECS

Fall 2009   John Kubiatowicz

Midterm I
October 19[th], 2009
CS162: Operating Systems and Systems Programming

Your Name:	
SID Number:	
Circle the letters of CS162 Login	First: a b c d e f g h I j k l m n o p q r s t u v w x y z Second: a b c d e f g h I j k l m n o p q r s t u v w x y z
Discussion Section:	

General Information:
This is a **closed book** exam. You are allowed 2 pages of notes (both sides). You may use a calculator. You have 3 hours to complete as much of the exam as possible. Make sure to read all of the questions first, as some of the questions are substantially more time consuming.

Write all of your answers directly on this paper. *Make your answers as concise as possible.* On programming questions, we will be looking for performance as well as correctness, so think through your answers carefully. If there is something about the questions that you believe is open to interpretation, please ask us about it!

Problem	Possible	Score
1	20	
2	18	
3	24	
4	20	
5	18	
Total	100	

[This page left for π]

3.14159265358979323846264338327950288419716939937510582097494459230781640628620899

Problem 1: True/False [20 pts]

Please *EXPLAIN* your answer in TWO SENTENCES OR LESS (Answers longer than this may not get credit!). Also, answers without an explanation *GET NO CREDIT.*

Problem 1a[2pts]: Apple was the first company to develop mice and overlapping windows.

True / False

Explain:

Problem 1b[2pts]: A direct mapped cache can sometimes have a higher hit rate than a fully associative cache with an LRU replacement policy (on the same reference pattern).

True / False

Explain:

Problem 1c[2pts]: Threads within the same process share the same heap and stack.

True / False

Explain:

Problem 1d[2pts]: A microkernel-style operating system uses multiple address spaces inside the operating system – with components such as the file system, network stack, and device drivers all running at user level.

True / False

Explain:

Problem 1e[2pts]: An operating system that implements on-demand paging on a machine with software TLB miss handling (such as MIPS) must use an inverted page table.

True / False

Explain:

Problem 1f[2pts]: If the banker's algorithm finds that it's safe to allocate a resource to an existing thread, then all threads will eventually complete.

True / False

Explain:

Problem 1g[2pts]: The Nachos operating system uses Mesa-style condition variables for all synchronization.

True / False

Explain:

Problem 1h[2pts]: The lottery scheduler prevents CPU starvation by assigning at least one ticket to each scheduled thread.

True / False

Explain:

Problem 1i[2pts]: Multicore chips (i.e. processor chips with more than one CPU on them) are only here for the short term (next few years) until the transistor feature size reaches 10nm.

True / False

Explain:

Problem 1j[2pts]: Thread pools are a useful tool to help prevent the "Slashdot" effect from crashing Web servers.

True / False

Explain:

Problem 2: Synchronization [18 pts]

Problem 2a[3pts]: Consider a Hash table with the following interface:

```
1.  public class HashTable {
2.     public void Put(int Key, int Value) {}
3.     public int Get(int Key) {}      // Return 0 if no previous Put() on Key
4.     public void Remove(int Key) {}  // No-op if no previous Put() on Key
5.  }
```

Assume that `Get()` must return the valid `Value` for any `Key` that has been written by previous `Put()` methods. If a `Put()` method on a given `Key` is happening at the same time as a `Get()` method, then the `Get()` method may return an earlier `Value`. In our attempt to make the HashTable threadsafe (i.e. usable by multiple threads at the same time), we might decided to make all three methods "synchronized" methods (e.g. Java synchronized statements). Would this have negative performance implications? Explain carefully (fully justify your answer; if the answer is "yes", explain why you think you could get better threadsafe performance. If the answer is "no", explain why this is the best threadsafe performance you could expect):

Problem 2b[2pts]: Explain the difference in behavior between `Semaphore.V()` and `CondVar.signal()` when no threads are waiting in the corresponding semaphore or condition variable:

Problem 2c[3pts]: Explain how **Nachos** is able to implement the correct semantics for `CondVar.signal()` using `Semaphore.V()`. Be explicit and make sure to explain why the different of (2b) is not an issue here.

Problem 2d[2pts]: Give two reasons why this is a bad implementation for a lock:

```
lock.acquire() { disable interrupts; }
lock.release() { enable interrupts; }
```

Problem 2e[8pts]: Suppose that we want a finite synchronized FIFO queue that can handle multiple simultaneous `enqueue()` and `dequeue()` operations. Assume that `enqueue()` and `dequeue()` are not allowed to busywait (but rather must sleep) when the queue is either full or empty respectively. Here is a sketch of an implementation utilizing a circular buffer:

```
1. static final int QUEUESIZE=100;
2. class FIFOQueue {
3.    Lock FIFOLock=new Lock();    // Methods acquire() and release()
4.    CondVar CV1=new CondVar(FIFOLock);// Methods wait(),signal(),broadcast()
5.    CondVar CV2=new CondVar(FIFOLock);// Methods wait(),signal(),broadcast()
6.    Object FIFO[QUEUESIZE];      // Finite circular queue of Objects
7.    int head = 0, tail = 0;      // Start out "empty"
8.
9.    void Enqueue(Object newobject)  {
10.       /* Enqueue Method. Spin until can enqueue */
11.    }
12.    Object Dequeue()  {
13.       /* Dequeue Method. Spin until can dequeue */
14.    }
15. }
```

Implement the `Enqueue()` and `Dequeue()` methods using monitor synchronization. You should enqueue using the tail variable and dequeue at the head. Remember that spin waiting is not allowed. You should have no more than 10 lines for each method. *Hint: make sure to account for wrapping of head and tail pointers and assume Mesa scheduling of the monitors.*

```
        void Enqueue(Object newobject)  {
```

```
        }
        Object Dequeue()  {
```

```
        }
```

[This page intentionally left blank]

Problem 3: Deadlock and the Cephalopod Banquet [24pts]

Problem 3a[4pts]: Name and explain the four conditions for deadlock**:**

Problem 3b[2pts]: Suppose that we utilize the Banker's algorithm to determine whether or not to grant resource requests to threads. The job of the Banker's algorithm is to keep the system in a "SAFE" state. It denies resource requests by putting the requesting thread to sleep if granting the request would cause the system to enter an "UNSAFE" state, waking it only when the request could be granted safely. What is a SAFE state?

Problem 3c[3pts]: Explain how the Banker's algorithm prevents deadlock by removing one or more of the conditions of deadlock from (3a). Be explicit.

Problem 3d[4pts]: Suppose that we wish to evaluate the current state of the system and declare whether or not it is in a SAFE state. In order to do this, we will need to keep explicit track of the resources in the system. In particular, if there were only two types of resource, we could describe the state of the system with the following data structures:

```
class FreeResources {
   int FreeResA, FreeResB;// Number of copies of resource that are free
}
/* Per-thread descriptor of thread resources */
class ThreadResources {
   int MaxNeededA, MaxNeededB;   // Max number copies of resource needed
   int CurHeldA, CurHeldB;       // Current number resources held
}
ThreadResources[] ThreadRes;
```

Assume that FreeRes and ThreadRes have been initialized to reflect the current state of the system. Here is a sketch for how we could check for safety:

```
1.  boolean IsSAFE(FreeResources FreeRes, ThreadResources[] ThreadRes) {
2.     Int FreeA = FreeRes.FreeResA, FreeB = FreeRes.FreeResB;
3.     boolean[] ThreadFinished = new boolean[ThreadRes.length];
4.     int RemainingThreads = ThreadRes.length;
5.     boolean finished = false;
6.     while (!finished) {
7.        finished = true;
8.        for (int i = 0; i < ThreadRes.length; i++) {
9.           if (!ThreadFinished[i]) {
10.          /* Missing Code */
11.          }
12.       }
13.    }
14.    return (RemainingThreads == 0);  /* SAFE if no threads left */
15. }
```

Provide the missing Code at line 10. **This code should have no more than 8 lines and should not alter the external variables (arguments).** *Hint: work through the threads that can complete.*

Code for Line 10:

The Cephalopod Diners Problem: Consider a large table with *identical* multi-even-armed cephalopods (e.g. octopuses). *In the center is a pile of forks and knives.* Before eating, each diner must have an equal number of forks and knives, one in each arm (e.g. if octopuses are eating, they would each need four forks and four knives). The creatures are so busy talking that they can only grab one utensil at a time. They also grab utensils in a random order until they have enough utensils to eat. After they finish eating, they return all of their utensils at once. Diners are implemented as threads that ask for utensils and return them when finished. Consider the following sketch for a CephTable class to implement the Cephalopod Diners problem using *monitor synchronization:*

```
1.  class DinerUtensils {
2.      public int forks,knives;          // utensils held by creature
3.  }                                     // clearly forks+Knives <= NumArms
4.  public class CephTable {
5.      Lock lock = new Lock();           // acquire(), release()
6.      CondVar CV = new CondVar(lock);//wait(),signal(),broadcast();
7.      public DinerUtensil[] Diners;     // Accounting: utensils for each diner
8.      int NumArms;                      // Number of arms for every diner
9.      int IdleForks, IdleKnives;        // Number of forks/knives on table
10.
11.     public CephTable(int NumDiners, int NumArms, int Forks, int Knives){
12.         Diners = new DinerUtensils[NumDiners]; // info about each Diner
13.         This.NumArms = NumArms;       // Number of arms per Diner
14.         This.IdleForks = Forks;       // Number Forks on table initially
15.         This.IdleKnives = Knives;     // Number Knives on table initially
16.     }
17.     public void GrabUtensil(int CephalopodID, boolean WantFork) {
18.         /* Try to grab a utensil from table */
19.     }
20.     public void DoneEating(int CephalopodID) {
21.         /* Return all chopsticks to pile */
22.         lock.acquire();
23.         IdleForks += Diners[CephalopodID].forks;
24.         IdleKnives += Diners[CephalopodID].knives;
25.         Diners[CephalopodID].forks = 0;
26.         Diners[CephalopodID].knives = 0;
27.         CV.broadcast();
28.         lock.release();
29.     }
30.     boolean CephCheck(int CephalopodID, int numforks, int numknives) {
31.         /* See if ok to give dinner numforks forks and numknives knives. */
32.     }
33. }
```

Problem 3e[3pts]: In its general form, the Banker's algorithm makes a decision about whether or not to allow an allocation request by making multiple passes through the set of resource holders (threads). See, for instance, the fact that there are two loops in (3d) to determine safety. Explain why a Banker's algorithm dedicated to the Cephalopod Diners problem, namely the CephCheck() routine, could operate with a single pass through the resources holders:

Problem 3f[4pts]: Implement the CephCheck method of the CephTable Object, namely fill in code for line 31 above. This method should implement the Banker's algorithm: return true if the given Cephalopod can be granted 'numforks' forks and 'numknives' knives without taking the system out of a SAFE state. Do not blindly implement the Banker's algorithm: this method only needs to have the same external behavior as the Banker's algorithm for this application. Note that this method is part of the CephTable Object and thus has access to local variables of that object. ***This code should not permanently alter the local variables of the CephTable Object (although it can do so temporarily).*** Do not worry about making this routing threadsafe; it will be called with a lock held. We will give full credit for a solution that takes a single pass through the diners, partial credit for a working solution, and no credit for a solution with more than 15 lines. *Hint: it is easier to first check the requesting Cephalopod, then the rest.*

Code for Line 31:

Problem 3g[4pts]: Implement the code for the GrabUtensil() routine, namely fill in code for line 18 above. Its behavior is that it should check to whether or not it is ok to grant the requested type of utensil to the caller and if not, sleep until it is ok. This code should call the CephCheck() routine as a subroutine and should be *threadsafe* namely, it should be able to deal with multiple threads accessing the state simultaneously. You should implement this routine as a monitor and assume Mesa scheduling. It is up to the Cephalopod to eat and subsequently call DoneEating(); you should not do that in GrabUtensil(). This routine can be written in 8 lines, but you can use up to 12:

Code for Line 18:

[This page intentionally left blank]

Problem 4: Virtual Memory [20 pts]

Consider a multi-level memory management scheme with the following format for virtual addresses:

Virtual Page # (10 bits)	Virtual Page # (10 bits)	Offset (12 bits)

Virtual addresses are translated into physical addresses of the following form:

Physical Page # (20 bits)	Offset (12 bits)

Page table entries (PTE) are 32 bits in the following format, *stored in big-endian form* in memory (i.e. the MSB is first byte in memory):

Physical Page # (20 bits)	OS Defined (3 bits)	0	Large Page	Dirty	Accessed	Nocache	Write Through	User	Writeable	Valid

Here, "Valid" means that a translation is valid, "Writeable" means that the page is writeable, "User" means that the page is accessible by the User (rather than only by the Kernel). *Note: the phrase "page table" in the following questions means the multi-level data structure that maps virtual addresses to physical addresses.*

Problem 4a[2pts]: How big is a page? Explain.

Problem 4b[2pts]: Suppose that we want an address space with one physical page at the top of the address space and one physical page at the bottom of the address space. How big would the page table be (in bytes)? Explain.

Problem 4c[2pts]: What is the maximum size of a page table (in bytes) for this scheme? Explain.

Problem 4d[2pts]: How big would each entry of a fully-associative TLB be for this management scheme? Explain.

Problem 4e[2pts]: Sketch the format of the page-table for the multi-level virtual memory management scheme of (4a). Illustrate the process of resolving an address as well as possible.

Problem 4f[10pts]: Assume the memory translation scheme from (4a). Use the Physical Memory table given on the next page to predict what will happen with the following load/store instructions. Assume that the base table pointer for the current ***user level process*** is `0x00200000.`

Addresses are virtual. The return value for a load is an 8-bit data value or an error, while the return value for a store is either "**ok**" or an error. Possible errors are: **invalid, read-only, kernel-only.**
Hint: Don't forget that Hexidecimal digits contain 4 bits!

Instruction	Result	Instruction	Result
Load [0x00001047]	0x50	Store [0x02001345]	
Store [0x00C07665]	ok	Load [0xFF80078F]	
Store [0x00C005FF]	ERROR: read-only	Load [0xFFFFF005]	
Load [0x00003012]		Test-And-Set [0xFFFFF006]	

Physical Memory [All Values are in Hexidecimal]

Address	+0	+1	+2	+3	+4	+5	+6	+7	+8	+9	+A	+B	+C	+D	+E	+F
00000000	0E	0F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D
00000010	1E	1F	20	21	22	23	24	25	26	27	28	29	2A	2B	2C	2D
….																
00001010	40	41	42	43	44	45	46	47	48	49	4A	4B	4C	4D	4E	4F
00001020	40	03	41	01	30	01	31	03	00	03	00	00	00	00	00	00
00001030	00	11	22	33	44	55	66	77	88	99	AA	BB	CC	DD	EE	FF
00001040	10	01	11	03	31	03	13	00	14	01	15	03	16	01	17	00
….																
00002030	10	01	11	00	12	03	67	03	11	03	00	00	00	00	00	00
00002040	02	20	03	30	04	40	05	50	01	60	03	70	08	80	09	90
00002050	10	00	31	01	10	03	31	01	12	03	30	00	10	00	10	01
….																
00004000	30	00	31	01	11	01	33	03	34	01	35	00	43	38	32	79
00004010	50	28	84	19	71	69	39	93	75	10	58	20	97	49	44	59
00004020	23	03	20	03	00	01	62	08	99	86	28	03	48	25	34	21
….																
00100000	00	00	10	65	00	00	20	67	00	00	30	00	00	00	40	07
00100010	00	00	50	03	00	00	00	00	00	00	00	00	00	00	00	00
…																
00103000	11	22	00	05	55	66	77	88	99	AA	BB	CC	DD	EE	FF	00
00103010	22	33	44	55	66	77	88	99	AA	BB	CC	DD	EE	FF	00	67
…																
001FE000	04	15	00	00	48	59	70	7B	8C	9D	AE	BF	D0	E1	F2	03
001FE010	10	15	00	67	10	15	10	67	10	15	20	67	10	15	30	67
…																
001FF000	00	00	00	00	00	00	00	65	00	00	10	67	00	00	00	00
001FF010	00	00	20	67	00	00	30	67	00	00	40	65	00	00	50	07
…																
001FFFF0	00	00	00	00	00	00	00	00	10	00	00	67	00	10	30	65
…																
00200000	00	10	00	07	00	10	10	07	00	10	20	07	00	10	30	07
00200010	00	10	40	07	00	10	50	07	00	10	60	07	00	10	70	07
00200020	00	10	00	07	00	00	00	00	00	00	00	00	00	00	00	00
…																
00200FF0	00	00	00	00	00	00	00	00	00	1F	E0	07	00	1F	F0	07
…																

Problem 5: Scheduling [18pts]

Problem 5a[2pts]: Give two ways in which to predict runtime in order to approximate SRTF:

Problem 5b[2pts]: What scheduling problem did the original Mars rover experience? What were the consequences of this problem?

Problem 5c[3pts]:

Five jobs are waiting to be run. Their expected running times are 10, 8, 3, 1, and X. In what order should they be run to minimize average completion time? State the scheduling algorithm that should be used AND the order in which the jobs should be run. *HINT: Your answer will explicitly depend on X.*
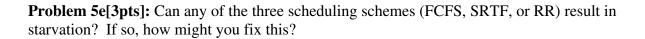
Problem 5d[5pts]:
Here is a table of processes and their associated arrival and running times.

Process ID	Arrival Time	CPU Running Time
Process 1	0	2
Process 2	1	6
Process 3	4	1
Process 4	7	4
Process 5	8	3

Show the scheduling order for these processes under 3 policies: First Come First Serve (FCFS), Shortest-Remaining-Time-First (SRTF), Round-Robin (RR) with timeslice quantum = 1. *Assume that context switch overhead is 0 and that new RR processes are added to the **head** of the queue and new FCFS processes are added to the **tail** of the queue.*

Time Slot	FCFS	SRTF	RR
0			
1			
2			
3			
4			
5			
6			
7			
8			
9			
10			
11			
12			
13			
14			
15			

Problem 5e[3pts]: Can any of the three scheduling schemes (FCFS, SRTF, or RR) result in starvation? If so, how might you fix this?

Problem 5f[3pts]: Explain why a chess program running against another program on the same machine might want to perform a lot of superfluous I/O operations.

[Scratch Page: Do not put answers here!]

[Scratch Page: Do not put answers here!]