University of California, Berkeley
College of Engineering
Computer Science Division — EECS

Fall 2005   John Kubiatowicz

Midterm I
October 12th, 2005
CS162: Operating Systems and Systems Programming

Your Name:	
SID Number:	
Discussion Section:	

General Information:
This is a **closed book** exam. You are allowed 1 page of **hand-written** notes (both sides). You have 3 hours to complete as much of the exam as possible. Make sure to read all of the questions first, as some of the questions are substantially more time consuming.

Write all of your answers directly on this paper. *Make your answers as concise as possible.* On programming questions, we will be looking for performance as well as correctness, so think through your answers carefully. If there is something about the questions that you believe is open to interpretation, please ask us about it!

Problem	Possible	Score
1	25	
2	17	
3	20	
4	18	
5	20	
Total		

[This page left for π]

3.14159265358979323846264338327950288419716939937510582097494944

Problem 1: Short Answer

Problem 1a[2pts]: Suppose a thread is running in a critical section of code, meaning that it has acquired all the locks through proper arbitration. Can it get context switched? Why or why not?

Problem 1b[3pts]: What are some of the hardware differences between kernel mode and user mode? Name at least three.

Problem 1c[3pts]: Name three ways in which the processor can transition from user mode to kernel mode. Can the user execute arbitrary code after transitioning?

Problem 1d[3pts]: What is a thread? What is a process? Describe how to create each of these.

Problem 1e[3pts]: Suppose that we have a two-level page translation scheme with 4K-byte pages and 4-byte page table entries (includes a valid bit, a couple permission bits, and a pointer to another page/table entry). What is the format of a 32-bit virtual address? Sketch out the format of a complete page table.

Problem 1f[2pts]: What needs to be saved and restored on a context switch between two threads in the same process? What if the two threads are in different processes? Be explicit.

Problem 1g[1pt]: What is a thread-join operation?

Problem 1h[3pts]: Name at least three ways in which context-switching can happen in a non-preemptive scheduler.

Problem 1i[3pts]: Name three ways in which processes on the same processor can communicate with one another. If any of the techniques you name require hardware support, explain.

Problem 1j[2pts]: What is an interrupt? What happens when an interrupt occurs? What is the function of an interrupt controller?

[This page intentionally left blank]

Problem 2: Monitors

Problem 2a[2pts]: What is a monitor?

Problem 2b[5pts]: Provide Java class definitions for each of the components of a monitor. Include specifications for local data and methods (names and arguments, not implementation); method names should adhere to the interface defined in class. Make sure to document each method by saying what it should do. Be brief! *Hint: you should have two different object classes.*

Problem 2c[2pts]: What is the difference between Mesa and Hoare scheduling for monitors?

Problem 2d[8pts]:   In parallel programs (one multi-threaded process), a common design methodology is to perform processing in sequential stages. All of the threads work independently during each stage, but they must synchronize at the end of each stage at a synchronization point called a *barrier.* If a thread reaches the barrier before all other threads have arrived, it waits. When all threads reach the barrier, they are notified and can begin execution on the next phase of the computation.

Example:

```
While (true) {
        Compute stuff;
        BARRIER();
        Read other threads results;
}
```

There are three complications to barriers. First, there is no master thread that controls the threads, waits for each of them to reach the barrier, and then tells them to restart. Instead, the threads must monitor themselves and determine when they should wait or proceed. Second, for many dynamic programs, the number of threads that will be created during the lifetime of the parallel program is unknown in advance, since a thread can spawn another thread, which will start in the same program stage as the thread that created it. Third, a thread may end before the barrier. In all cases, all threads must wait at the barrier for all other threads before anyone is allowed to proceed.

Provide the pseudo-code for a monitor class called **Barrier** that enables this style of barrier synchronization. Your solution must support creation of a new thread (an additional thread that needs to synchronize), termination of a thread (one less thread that needs to synchronize), waiting when a thread reaches the barrier early, and releasing waiting threads when the last thread reaches the barrier. *Implement your solution using monitors. You should never busy-wait; threads waiting at a barrier should be sleeping.*

Your   class   must   implement   the   following   three   methods: `threadCreated()`, `threadEnd()`, and `enterBarrier()`.   Think carefully about efficiency and avoid unnecessary looping. Feel free to utilize space on the next page.

[This page intentionally left blank]

[This page intentionally left blank]

Problem 3: Lock-Free Queue

An object such as a queue is considered "lock-free" if multiple processes can operate on this object simultaneously without requiring the use of locks, busy-waiting, or sleeping. In this problem, we are going to construct a lock-free FIFO queue using an atomic "swap" operation. This queue needs both an `Enqueue` and `Dequeue` method.

We are going to do this in a slightly different way than normally. Rather than `Head` and `Tail` pointers, we are going to have "PrevHead" and `Tail` pointers. `PrevHead` will point at the last object returned from the queue. Thus, we can find the head of the queue (for dequing). Here are the basic class definitions (assuming that only one thread accesses the queue at a time):

```
// Holding cell for an entry
class QueueEntry {
    QueueEntry next = null;
    Object stored;
    int taken = 0;

    QueueEntry(Object newobject) {
        stored = newobject;
    }
}

// The actual Queue (not yet lock free!)
class Queue {
    QueueEntry prevHead = new QueueEntry(null);
    QueueEntry tail = prevHead;

    void Enqueue(Object newobject)  {
        newEntry = new QueueEntry(newobject);
        tail.next = newEntry;
        tail = newEntry;
    }

    Object Dequeue() {
        QueueEntry nextEntry = prevHead.next;
        While ((nextEntry != null) && nextEntry.taken == 1) {
            nextEntry = nextEntry.next;
        }
        if (nextEntry == null) {
            return null;
        } else {
            nextEntry.taken = 1;
            prevHead = nextEntry;
            return nextEntry.stored;
        }
    }
}
```

Problem 3a[10pts]:

Suppose that we have an atomic swap instruction. This instruction takes a local variable (register) and a memory location and swaps the contents. Although Java doesn't contain pointers, we might describe this as follows:

```
Object AtomicSwap(variable addr, Object newValue) {
    object result = *addr;     // get old object stored in addr
    *addr = newValue;          // store new object into addr
    return result;             // Return old contents of addr
}
```

With this primitive, we might build a test-and-set lock as follows:

```
int mylock = 0;

// grab lock
while (AtomicSwap(mylock,1) == 1);
 // got lock
```

Rewrite code for Enqueue(), using the AtomicSwap() operation, such that it will work for any number of simultaneous Enqueue and Dequeue operations. You should never need to busy wait. **Do not use locking (i.e. don't use a test-and-set lock).** The solution is tricky but can be done in a few lines. We will be grading on conciseness. *Hint: during simultaneous insertions, objects may be temporarily disconnected from the queue (i.e. the set of entries reachable by* nextEntry = nextEntry.next *starting from* prevEntry); *but eventually the queue needs to be connected.*

Problem 3b[10pts]:
Rewrite code for Dequeue() such that it will work for any number of simultaneous threads working at once. **Again, do not use locking.** You should never need to busy wait. The solution is tricky but can be done by modifying a small number of lines. We will be grading on conciseness. *Hint: before grabbing a queue entry, make sure that you "take" it atomically.* You are striving for correctness. If you can explain why it is ok, you can allow prevHead to get temporarily out of date, but try to avoid letting prevHead get too far behind the actual head of the list.

Problem 4: Scheduling

Problem 4a[2pts]:
Six jobs are waiting to be run. Their expected running times are 10, 8, 6, 3, 1, and X. In what order should they be run to minimize average completion time? State the scheduling algorithm that should be used AND the order in which the jobs should be run. (Your answer will depend on X).

Problem 4b[4pts]:
Name and describe 4 different scheduling algorithms. What are the advantages and disadvantages of each?

Problem 4c[3pts]:

Here is a table of processes and their associated running times. *All of the processes arrive in numerical order at time 0.*

Process ID	CPU Running Time
Process 1	2
Process 2	6
Process 3	1
Process 4	4
Process 5	3

Show the scheduling order for these processes under 3 policies: First Come First Serve (FCFS), Shortest-Remaining-Time-First (SRTF), Round-Robin (RR) with timeslice quantum = 1

Time Slot	FCFS	SRTF	RR
0			
1			
2			
3			
4			
5			
6			
7			
8			
9			
10			
11			
12			
13			
14			
15			

Problem 4d[3pts]:
For each process in each schedule above, indicate the queue wait time and completion time (otherwise known as turnaround time, TRT). Note that wait time is the total time spend waiting in queue (all the time in which the task is not running):

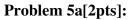| Scheduler	Process 1	Process 2	Process 3	Process 4	Process 5
FCFS wait					
FCFS TRT					
SRTF wait					
SRTF TRT					
RR wait					
RR TRT					

Problem 4e[3pts]:
Suppose that the context-switch time is 0.5 units. Assume that there is no context switch at time zero. Fill out this table again:

Scheduler	Process 1	Process 2	Process 3	Process 4	Process 5
FCFS wait					
FCFS TRT					
SRTF wait					
SRTF TRT					
RR wait					
RR TRT					

Problem 4f[3pts]:
The SRTF algorithm requires knowledge of the future. Why is that? Name two ways to approximate the information required to implement this algorithm.

[This page intentionally left blank]

[This page intentionally left blank]

Problem 5: Deadlock

Problem 5a[2pts]:
List the conditions for deadlock.

Problem 5b[3pts]:
What are three methods of dealing with the *possibility* of deadlock? Explain each method and its consequences to programmer of applications.

Problem 5c[3pts]:
Suppose that we have several processes and several resources. Name three ways of preventing deadlock between these processes

Problem 5d[5pts]:

Suppose that we have the following resources: A, B, C and threads T1, T2, T3, T4. The total number of each resource is:

Total		
A	B	C
12	9	12

Further, assume that the processes have the following maximum requirements and current allocations:

Thread ID	Current Allocation			Maximum		
	A	B	C	A	B	C
T1	2	1	3	4	3	4
T2	1	2	3	5	3	3
T3	5	4	3	6	4	3
T4	2	1	2	4	1	2

Is the system potentially deadlocked (i.e. unsafe)? In making this conclusion, you must show all your steps, intermediate matrices, etc.

Problem 5e[3pts]:
Assuming the allocations in (5d), suppose that T1 asks for 2 more copies of A. Can the system grant this or not? Explain.

Problem 5f[4pts]:
Assuming the allocations in (5d), what is the maximum number of additional copies of resources (A, B, and C) that T1 can be granted in a single request without risking deadlock? Explain.

[This page intentionally left blank]

[This page intentionally left blank]